



Input: frozen LLM (L decoder layers) · eviction policy · calibration subset (search) · full task data (selection)

Stage 1 Explore



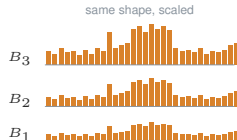
Unconstrained BO over per-layer budgets, trading average budget against calibration score.

Stage 2 Select



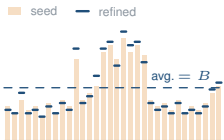
Re-evaluate the Pareto front on **full data** and keep the best allocation.

Stage 3 Transfer



Rescale the winner to every target budget with **no new search**.

Stage 4 Refine



BO at a fixed average budget, seeded with the winner; candidates are re-checked on full data.

Once per (task, eviction policy) – reused at every budget

Per target budget